

DEEP LEARNING – PART 1

Verena Kaynig-Fittkau
Ray (Thouis) Jones

ComputeFest 2016

Check This (Repository) Out

- https://github.com/vkaynig/IACS_ComputeFest_DeepLearning
- Or go to github and search for IACS_ComputeFest

What We Cover Today:

- Basic theory of deep neural networks (DNN)
- Introduction to Theano
- Training deep neural networks with Theano
- Some useful tips and tricks for training DNNs

Workshop Goal

- Flatten the learning curve for the Theano tutorials
- <http://deeplearning.net/tutorial/>

Deep Learning




Why do you want to know about deep learning?

Motivation

- It works!
 - State of the art in machine learning
 - Google, Facebook, Twitter, Microsoft are all using it.
-
- It is fun!
 - Need to know what you are doing to do it well.

Google Trends




Google Brain - 2012



16 000

What it learned




<http://www.nytimes.com/2012/06/26/technology/in-a-big-network-of-computers-evidence-of-machine-learning.html?pagewanted=all>

Manual Feature Design

Yes

No

classification




hand designed
features



input image
patch


Learned Feature Hierarchy



Deep Learning Techniques

- Artificial neural network
 - Introduced in the 60s
 - [Perceptron in the old days](#)
- Convolutional neural network
 - Introduced in the 80s
- Recurrent neural network
 - Introduced in the 80s

What Changed -Computational Power




© Google

What Changed – Data Size



Scaling with Data Size



I don't Have a Cluster at Home

GOOGLE BRAIN

1,000 CPU Servers
2,000 CPUs • 16,000 cores

600 kWatts
\$5,000,000

300X energy efficiency
400X lower cost
Fits under a desk



1 Titan Z-Accelerated Server
3 Titan Zs • 17,280 cores


2 kWatts
\$12,000

Deep Learning




What is deep learning?


What – Neural Network Architecture



Artificial Neuron





Artificial Neuron




$$s(b + w^T x)$$

Step vs Sigmoid Activation




$$s(x) = \frac{1}{1 + e^{-cx}}$$

Multi-Layer Perceptron



$$s(b^{(1)} + W^{(1)}x)$$

Multi-Layer Perceptron



$$f(x) = G(b^{(2)} + W^{(2)} (s(b^{(1)} + W^{(1)}x)))$$

G : logistic function, softmax for multiclass

Logistic Regression (Multiclass)

- Probabilistic classifier

$$\begin{aligned} P(Y = i|x, W, b) &= \text{softmax}_i(Wx + b) \\ &= \frac{e^{W_i x + b_i}}{\sum_j e^{W_j x + b_j}} \end{aligned}$$

$$y_{pred} = \operatorname{argmax}_i P(Y = i|x, W, b)$$

How - Optimization

- As so many machine learning methods, also DNNs are about optimizing a loss function:

$$\arg \min_{\boldsymbol{\theta}} \frac{1}{T} \sum_t l(f(\mathbf{x}^{(t)}; \boldsymbol{\theta}), y^{(t)}) + \lambda \Omega(\boldsymbol{\theta})$$

Diagram illustrating the optimization objective function with annotations:

- $\boldsymbol{\theta}$: model parameter (indicated by an upward arrow from "model parameter")
- $\mathbf{x}^{(t)}$: training samples (indicated by a downward arrow from "training samples")
- $y^{(t)}$: training labels (indicated by an upward arrow from "training labels")
- $\Omega(\boldsymbol{\theta})$: regularizer (indicated by a downward arrow from "regularizer")

Loss Function

- Classification error would be good, but it's not smooth
- Need to find smooth proxy
- Last layer is logistic regression
- The whole network estimates class probabilities

$$P(Y = i|x, W, b) = \text{softmax}_i(Wx + b)$$

- Maximize the correct class assignment probabilities in the training set

Minimize Negative Log Likelihood

- Maximize the likelihood of the data set given the model parameters
- Take the log to simplify for numerical stability and math simplicity
- Take negative log likelihood to cast into minimization

training data
↓

$$\mathcal{L}(\theta, \mathcal{D}) = \sum_{i=0}^{|\mathcal{D}|} \log P(Y = y^{(i)} | x^{(i)}, \theta)$$


↑
Parameter θ , \mathbf{W} , \mathbf{b}

Stochastic Gradient Decent (SGD)

- Take small steps downward on the surface of the loss function
- For each training sample: $\Delta = -\nabla_{\theta} l(f(\mathbf{x}^{(t)}; \theta), y^{(t)})$
 - Estimate gradient
 - Update Parameters $\theta \leftarrow \theta + \alpha \Delta$

\uparrow
learning rate
- We minimize negative log likelihood
- Need to compute gradient – Ugh
- Theano does it for us!!

Stochastic Gradient Decent



Libraries

- Theano
- Torch
- Caffe
- TensorFlow
- ...


Theano

- Full disclosure: My favorite
- Python
- Transparent GPU integration
- Symbolical Graphs
- Auto-gradient
- Low level – in a good way!
- If you want high-level on top:
 - Pylearn2
 - Keras, Lasagne, Blocks
 - ...

theano

Torch

- Lua (and no Python interface)
- Very fast convolutions
- Used by Google Deep Mind, Facebook AI, IBM
- Layer instead of graph based



Caffe

- C++ based
- Higher abstraction than Theano or Torch
- Good for training standard models
- Model zoo for pre-trained models

Tensorflow

- Symbolic graph and auto-gradient
- Python interface
- Visualization tools
- Some performance issues regarding speed and memory



We are going to use Theano today

- Python library for efficient evaluation of mathematical expressions
- **tight integration with NumPy** – Use *numpy.ndarray* in Theano-compiled functions.
- **transparent use of a GPU** – Perform data-intensive calculations up to 140x faster than with CPU.(float32 only)
- **efficient symbolic differentiation** – Theano does your derivatives for function with one or many inputs.

Free and Open Source

- F. Bastien, P. Lamblin, R. Pascanu, J. Bergstra, I. Goodfellow, A. Bergeron, N. Bouchard, D. Warde-Farley and Y. Bengio
- Giving credit:
<http://deeplearning.net/software/theano/citation.html#citation>

Exercise: clone github repository

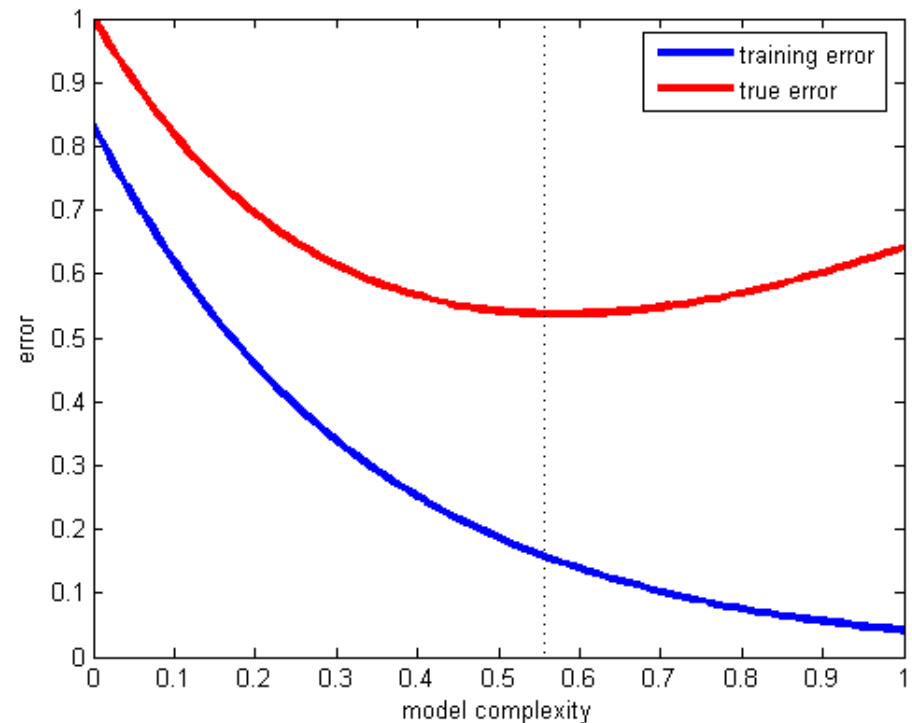
- `git clone --recursive`
https://github.com/vkaynig/IACS_ComputeFest_DeepLearning.git

LET'S GET SOME
EXPERIENCE

TIPS AND TRICKS

Number of Layers / Size of Layers

- If data is unlimited larger and deeper should be better
- Larger networks can over fit more easily
- Take computational cost into account



Learning Rate

- One of the most important parameters
- If network diverges most probably learning rate is too large
- Smaller works better
- Can slowly decay over time
- Can have one learning rate per layer

Other tips for SGD:

<http://leon.bottou.org/publications/pdf/tricks-2012.pdf>

Momentum

- Helps to escape local minima
- Crucial to achieve high performance

$$v_{t+1} = \mu v_t - \varepsilon \nabla f(\theta_t)$$

$$\theta_{t+1} = \theta_t + v_{t+1}$$

More about Momentum:

<http://www.jmlr.org/proceedings/papers/v28/sutskever13.pdf>

Convergence

- Monitor validation error
- Stop when it doesn't improve within n iterations
- If learning rate decays you might want to adjust number of iterations

Initialization of W

- Need randomization to break symmetry
- Bad initializations are untrainable
- Most heuristics depend on the number of input (and output) units
- Sometimes W is rescaled during training
 - Weight decay (L2 regularization)
 - Normalization

Data Augmentation

- Exploit invariances of the data
- Rotation, translation
- Nonlinear transformation
- Adding Noise

Type ◆	Classifier ◆
Neural network	6-layer NN 784-2500-2000-1500-1000-500-10 (on GPU), with elastic distortions
Convolutional neural network	Committee of 35 conv. net, 1-20-P-40-P-150-10, with elastic distortions

Preprocessing ◆	Error rate (%) ◆
None	0.35 ^[17]
Width normalizations	0.23 ^[8]

Data Normalization

- We have seen std and mean normalization
- Whitening
 - Neighbored pixels often are redundant
 - Remove correlation between features

More about preprocessing:

http://deeplearning.stanford.edu/wiki/index.php/Data_Preprocessing

Non-Linear Activation Function


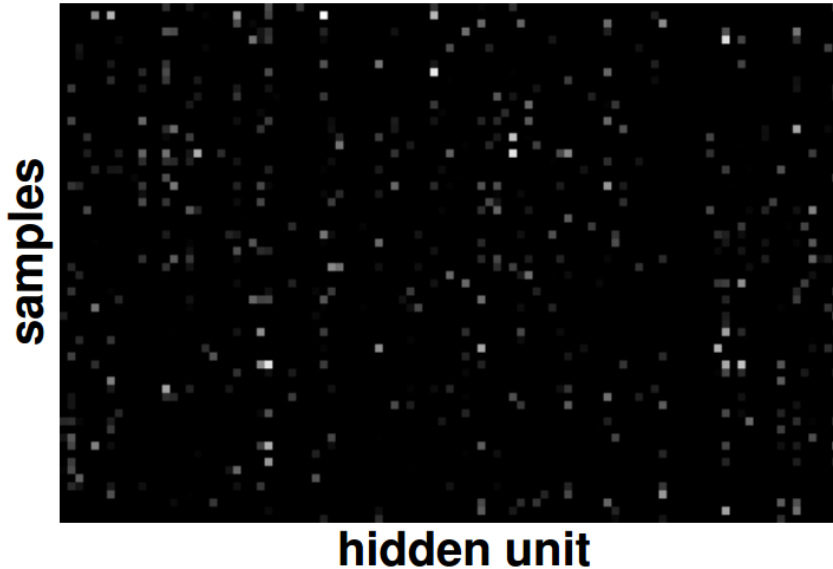
- Sigmoid
 - Traditional choice
- Tanh
 - Symmetric around the origin
 - Better gradient propagation than Sigmoid
- Rectified Linear (discussed tomorrow)

L1 and L2 Regularization

- Most pictures of nice filters involve some regularization
- L2 regularization corresponds to weight decay
- L2 and early stopping have similar effects
- L1 leads to sparsity
- Might not be needed anymore (more data, dropout)

Monitoring Training


- Monitor training and validation performance
- Can monitor hidden units
- Good: Uncorrelated and high variance



Autoencoder

- This is what Google used for their Google brain
- Basically just a MLP
- Output size is equal to input size
- Popular for pre-training a network on unlabeled data

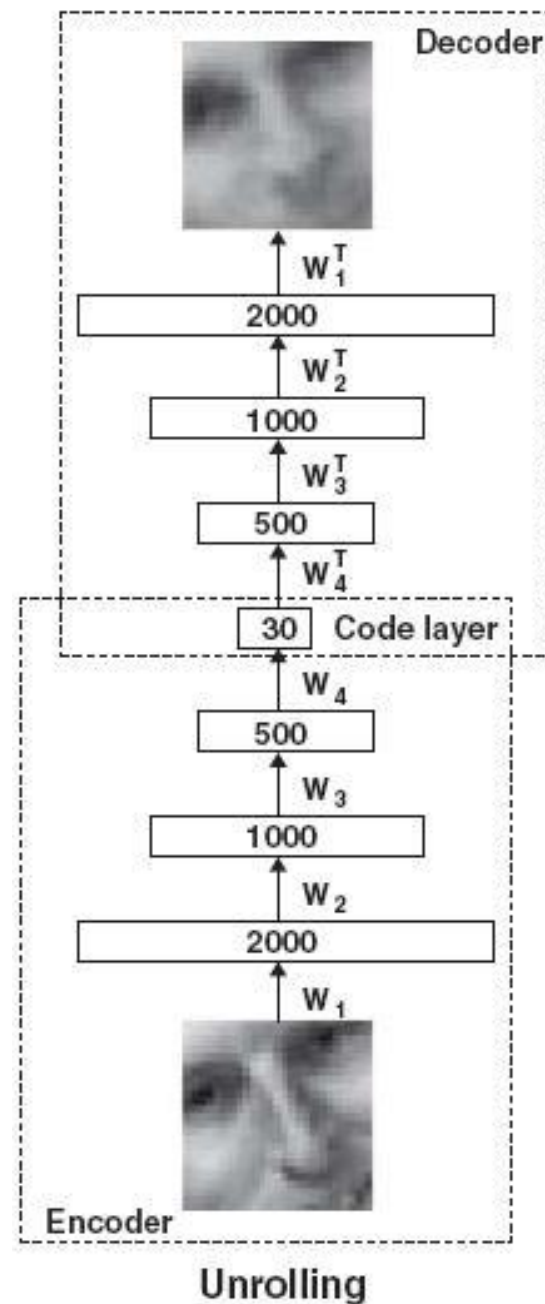
Autoencoder



- Reconstruct image from learned low dimensional code
- Weights are tied
- Learned features are often useful for classification

Deep Autoencoder

- Each Layer trained by itself
- Higher layers build on top of lower layers
- Can do fine tuning in the end



Further Resources

- More about theory:
 - Yoshua Bengio's book: <http://www.iro.umontreal.ca/~bengioy/dlbook/>
 - Deep learning reading list: <http://deeplearning.net/reading-list/>
- More about Theano:
 - <http://deeplearning.net/software/theano/>
 - <http://deeplearning.net/tutorial/>